

# CIS4301 Notes

Ryan Roden-Corrent

## 1 Functions of a Database

- **Create** - Add new records
- **Query** - Search for records meeting a description
- **Store** - Hold data in records
- **Durability** - data remains available and uncorrupted (backups)  
Example: Oracle provides prompt maintenance if a server you bought from them goes down
- **Access** - Manageable permissions. Control who can see/edit data

## 2 History of Databases

### 1870s

Charles Babbage and Ada Lovelace create the difference engine, and later, the analytic engine (mechanical computers)

### 1876

Invention of Dewey Decimal system: a general cataloging system for libraries that is still in use today.

### 1962

Bachman: integrated data store

- **network data model**: a tree with no loops
- **graph data model**: loops allowed

models are specific to the data being represented

### 1960s

**Disks** (random access) replace **tapes** (serial access). Now one can access data without reading the preceding data.

At this time, people interested in databases are mostly programmers.

IBM Builds **IMS** (Information Management System) to track Apollo.

**CODASYL**: Committee standardized network data model

## 1970s

Edward Codd writes 'A relational model of data for large shared banks'. It conflicts with the design of the current most popular databases. Eventually, IBM agrees to dedicate a research team to Codd's idea; this team produces SQUARE

## 1977

Another IBM team develops **Query By Example (QBE)**.

A group of CIA employees branch off to found **Oracle** and create the first commercial implementation of **SEQUEL**.

**SEQUEL** stands for **Structured English Query Language** but is abbreviated to **SQL** to avoid a copyright dispute.

## 1990s

Research into object oriented database implementations

## 2000s

Web Boom: Many services suddenly need to store massive amounts of data.

## 3 NOSQL

1970 No SQL - SQL doesn't exist

1980 SQL invented, standard for databases

1990 No, SQL - If you want to do database stuff, use SQL

2000 Not Only SQL - There are other things out there

2010 No, really, use SQL

## 4 Is a filesystem a database?

**Create** can create and remove records (touch, rm, mkdir, rmdir)

**Query** can search records (find, windows search bar)

**Store** files can store content that can be read back

**Durability** debateable

**Access** issues arise if two sources try to edit the same file

Which of these services are actually part of the filesystem, and which are really provided by the operating system?

## 5 Applications Of Databases

**Bank** Store Account information

**Social Media Site** User info, posts, need to find users/posts by query

**Advertising** Store and utilize user data for targeted advertising

**Music Library** Query music by artist, album, or other property

**Employee Records** Name, salary, ect.

### 5.1 Early Database Adopters

Early adopters of database systems were banks, airlines, telecom companies, and companies interested in tracking employee data.

### 5.2 Large Database Users

1. Google/Youtube
2. NSA
3. Amazon
4. Yahoo
5. Yahoo
6. Facebook
7. Sprint
8. AT&T
9. Apple
10. National Energy Research Scientific Computing (NERSC)
11. Lexis Nexis

All of these companies keep track of **user-generated content**

## 6 An overview of DMBS

**DDL:** Data Definition Language

CREATE, ALTER, DROP, TRUNCATE, COMMENT, RENAME

**DML:** Data Manipulation Language

SELECT, INSERT, UPDATE, DELETE


### 6.1 Flow of a query

*Query*  $\rightarrow$  *QueryParser*  $\rightarrow$  *QueryPlan*  $\rightarrow$  *QueryPreprocessor*  $\rightarrow$  *QueryOptimizer*  $\rightarrow$  *QueryPlan*

SQL is declarative - say what you want, and an optimal search method is chosen for you


## 6.2 Query Structure

$f_2(f_1(\text{Table}_1, \text{Table}_2)), \text{Table}_3$



The query optimizer may collapse this

$f_3(\text{Table}_1, \text{Table}_3)$



## 6.3 Parts of a Database

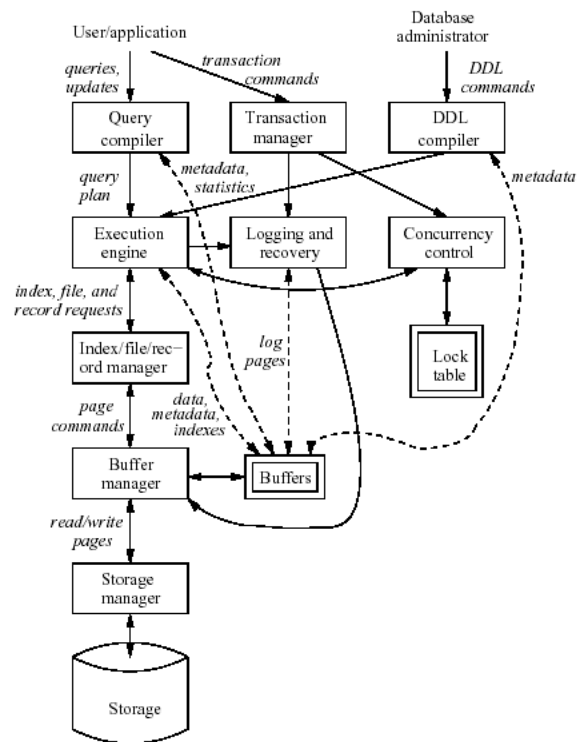


Figure 1: Database Flow Diagram

**Query Compiler** Form query

**Index/File/Record Manager** Can speed up queries using index

**Buffer Manager** Decide whether to cache in RAM (much faster than disk!)

**Storage Manager** Provide access to data on disk in *pages*

**Transaction Manager** Process transactions (see next section)

## 7 Transactions

Transactions are operations that must be executed reliably and consistently. Order is important, and transactions should not interfere with each other. In **multi-threaded** applications, race-conditions and deadlocks can occur while trying to execute transactions. Deadlock resolution must occur. Transactions are atomic - they need to happen all at once or not at all. The transaction manager must initiate **rollback** (undo) in case of an error. Transactions are run in isolation from other database operations.

### 7.1 ACID

**Atomicity** transaction happens just once or not at all

**Consistency** predictable ordering, same result every time

**Isolation** no interference between transactions

**Durability** no data missing after transaction

### 7.2 Transaction Example

```
cg == 1
T_1
BEGIN
1/13 | cg | +5
1/14 | cg | -1
END;
cg == 5
```

```
T_2
BEGIN
1/13 | cg | -5
1/13 | cg | +1
END;
```

End result should be the same regardless of order of execution

```
T_1; T_2 == 1
T_1; T_2 == 1
```

## 8 Relational Data Model

1. Physical Model
2. Operations on data
3. Constraints on data

## 8.1 Physical Model Example

---

```
public class Animal{\n{\n    public int Eyes;\n    public String birthDate;\n    public double Weight;\n\n    public void beCute();\n}\n\nAnimal a = new Animal();\n[header] [integer] [birthDate...] [double] [beCute]
```

---

Constraints:

- birthDate can't be null
- Eyes can't be negative

## 8.2 Relational Model

Separate data from operations.

**Relations** no more than 2

**Attributes** order matters, **atomic** (int,double,String)

**Tuples**

**Domain**

**Keys** uniquely define attribute instance

## 8.3 Relational Model Example

---

```
Animal(eyes:integer, birthDate:String, weight:double)\n(0001, 2, "November 2", 12.0) //cannot have 2 of these, must be unique\n(0002, 3, "November 2", 12.0)\n\n//don't have to store entire animal, could just store certain columns\nvar a[] = [2,3]
```

---

## 8.4 Thinking with Relational Models

Don't think about how data will be stored, just how it will be used.

Entity Relational Model:

**Entity/Sets**

**Relationships**

**Attributes**

## 8.5 Constraints

**many-many**

**one-one**

**one-many/many-one** : a pet only has one owner, an owner may have many pets