

CIS4301 Notes

Ryan Roden-Corrent

February 5, 2014

1 Relational Algebra

See <http://cise.ufl.edu/class/cis4301sp14/slides/ra.ppt> for class slides on this. These notes are mostly a condensed version of the slides, the slides contain some nice table images to help you visualize the operations.

1.1 What is Relational Algebra?

Operators most common actions you execute on relations

Operands relations or variables that represent relations

1.2 Core Relational Algebra

Union, Intersection, Difference most common actions you execute on relations

Selection picking certain rows

Projection picking certain columns

Products/Joins compositions of relations

Renaming of relations and attributes

1.2.1 Selection

$R1 := \sigma_c(R2)$

C is a condition that refers to attributes of R2

R1 is all tuples of R2 that satisfy C

Relation Sells:			$JoeMenu := \sigma_{bar="Joe's"}(Sells) :$		
bar	beer	price	bar	beer	price
Joe's	Bud	2.50	Joe's	Bud	2.50
Joe's	Miller	2.75	Joe's	Miller	2.75
Sue's	Bud	2.50	Sue's	Bud	2.50
Sue's	Miller	3.00	Sue's	Miller	3.00

1.2.2 Projection

$R1 := \pi_L(R2)$

L is a list of attributes from R2's schema

R1 contains only the attributes of R2 listed in L (in the order they are listed)

Set operation: removes duplicate tuples

1.2.3 Extended Projection

$R1 := \pi_L(R2)$

Like projection, but L can contain arbitrary expressions involving attributes. Example: $\pi_{A+B->C,A,A}(R)$ will create a new table where the first column, C, is the sum of A and B, and the next two columns A1 and A2 are copies of the original A.

1.2.4 Product

$R3 := R1 \times R2$

- pair each tuple t1 of R1 with each tuple t2 of R2
- Concatenation $t_1 t_2$ is a tuple of R_3
- Schema of R3 is the attributes of R1 and then R2, in order
- Beware attribute A of same name in R1 and R2: use R1.A and R2.A

Example: $R3 := R1 \times R2$

Duplicate column B, so use R1.B and R2.B to differentiate.

1.2.5 Theta Join

$R3 := R1 \bowtie_C R2$ Equivalent to taking the product $R1 \times R2$ and applying σ_C to the result.

$R \bowtie_\theta S \equiv \sigma_\theta(R \times S)$

C can be any boolean-values condition.

1.2.6 Natural Join

Remove duplicate columns, only return columns that naturally combine. $R3 := R1 \bowtie R2$

1.2.7 Renaming

Gives a new schema to a relation.

$R1 := \rho_{R1(A_1, \dots, A_n)}(R2)$ or $R1(A_1, \dots, A_n) := R2$ (simplified notation)

R1 is a relation with the same tuples as R2 but the attributes A_1, \dots, A_n .

1.3 Building Complex Expressions

Combine operators with parentheses and precedence rules.

Three notations:

1. Sequences of assignment statements
2. Expressions with several operators
3. Expression trees

1.3.1 Operator Precedence

1. $[\sigma, \pi, \rho]$ (highest)
2. $[X, \bowtie]$
3. \cap
4. \cup

1.3.2 Expression Trees

Leaves are operands (variables or constant relations).

Interior nodes are operators applied to children.

Example:

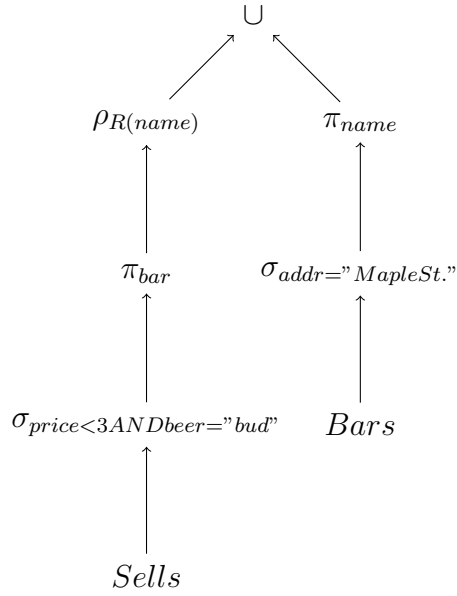


Figure 1: Find the names of all bars that are either on maple street or sell Bud for less than \$3

1.4 Relational Algebra on Bags

A bag is like a set, but duplicate elements are allowed. SQL is a bag language. Operations like projection are more efficient on bags than on sets.

1.4.1 Bag Union

Just add elements together, including duplicates.

$$\{1, 2, 1\} \cup \{1, 1, 2, 3, 1\} = \{1, 1, 1, 1, 1, 2, 2, 3\}$$

1.4.2 Bag Intersection

Minimum number of duplicates in resulting set.

$$\{1, 2, 1, 1\} \cap \{1, 2, 1, 3\} = \{1, 1, 2\}$$

1.4.3 Bag Difference

Result contains all tuples in first relation that aren't in the second.

$$\{1, 2, 1, 1\} - \{1, 2, 3\} = 1, 1$$

1.4.4 Bag Laws != Set Laws

Set union is **idempotent**, (result does not change if applied multiple times.)

For a bag union, if x appears n times in S , then it appears $2n$ times in union.

2 Why SQL?

Say what to do rather than how to do it.

SELECT desired attributes

FROM one or more tables

WHERE some condition holds

2.1 Example:

Names of beers made by Anheuser-Busch:

```
SELECT name
FROM Beers
WHERE manf = 'Anheuser-Busch'
```

name	manf
Bud	Anheuser-Busch

* indicates "all attributes in relation.

```
SELECT name
FROM Beers
WHERE manf = 'Anheuser-Busch'
```

Try adding things to SELECT *

Rename *name* to *beer* using **AS**

```
SELECT name AS beer, manf
FROM Beers
WHERE manf = 'Anheuser-Busch'
```

```
SELECT bar, ber,
price*114 AS PriceInYen
FROM Beers
```

Using **Likes**(*drinker*, *beer*): Put 'likes Bud' in attribute whoLikesBud

```
SELECT drinker
  'likes Bud' AS whoLikesBud
FROM Likes
WHERE beer = 'Bud'
```

2.2 Information Integration

Each **Data warehouse** manages several subdatabases that each have multiple tables.

Example: Each bar has its own relation Menu(beer,price) tat we need to integrate.

```
SELECT 'Joe''s Bar', beer, price
FROM Menu;
```

Do this for each bar to create a consistent schema, then take the Union.

Note: the repeated single quote is used to escape the nested quote.

2.3 Complex Conditions

Using Sells(bar, beer, price), find the price Joe charges for Bud.

```
SELECT price
FROM Sells
WHERE bar = 'Joe''s Bar' AND beer = 'Bud';
```

2.4 Patterns

Usable on **VARCHAR**. Useable on **DATE**, **INT**? Maybe.

```
<Attribute> LIKE <pattern>
<Attribute> iLIKE <pattern>    --case insensitive like
<Attribute> NOT LIKE <pattern>
```

% means any string

_ means any character

Example: choose numbers with two digits

```
FROM NUMBERS
WHERE number LIKE "__"
```

Example: choose numbers with at least two digits

```
FROM NUMBERS
WHERE number LIKE "__%"
```

Find drinkers with phone number starting with 555.

```
SELECT name
FROM Drinkers
WHERE phone LIKE '%555- _ _ _ _';
```

2.5 NULL Values

NULL could mean:

Missing Value we know it exists but don't have value

Inapplicable e.g. spouse of unmarried person

NULLs can be tricky, try to avoid if possible.

2.5.1 3-valued logic

TRUE = 1, FALSE = 0, UNKNOWN = 1/2

AND = MIN, OR = MAX, NOT(x) = 1-x

Ex:

TRUE AND (FALSE OR NOT(UNKNOWN)) = MIN(1, MAX(0, (1-1/2)))

Start on innermost paren and change truth values to numeric values. E.g. NOT(UNKNOWN) = (1-1/2)

2.5.2 Surprising Example

bar	beer	price
Joes	Bud	NULL

```
SELECT bar
FROM Sells
WHERE price < 2.00 OR price >= 2.00;
```

The comparisons become UNKNOWN because they are undefined for NULL.
Result is

NULL OR NULL
Max(1/2, 1/2) = 1/2

Expected to get everything, got nothing

2.6 Multirelation Queries

Distinguish between attributes of same name by "relation.attribute"

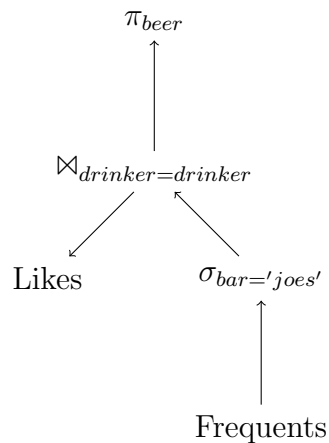
2.6.1 Examples

Using `Likes(drinker,beer)`, `Frequents(drinker, bar)`, find the beers liked by at least one person who frequents Joe's.

```
SELECT beer
FROM Likes, Frequents
WHERE Bar = 'Joes' AND
      Frequents.drinker =
      Likes.drinker;
```

A theta join $\bowtie_{=}$ is used to equate the drinker attribute from each table. Can be done in any order.

For example:



3 Multirelational Queries

3.1 Joining Two Relations

Find the beers liked by at least one person who frequents Joe's

```
SELECT beer
FROM Likes, Frequents
WHERE bar = 'Joes' AND
      Frequents.drinker =
      Likes.drinker;
```

Same as a $\bowtie_{Frequents.drinker=Likes.drinker}$

3.2 Self-Join

Get all pairs of beers that have same manufacturer and a different name.

```

SELECT b1.name, b2.name
FROM Beers b1, Beers b2 --avoid naming conflicts
WHERE b1.manf = b2.manf AND
      b1.name < b2.name;

```

Why use $\boxed{<}$ instead of $\boxed{<>}$? $<$ will order the results and get rid of duplicates.

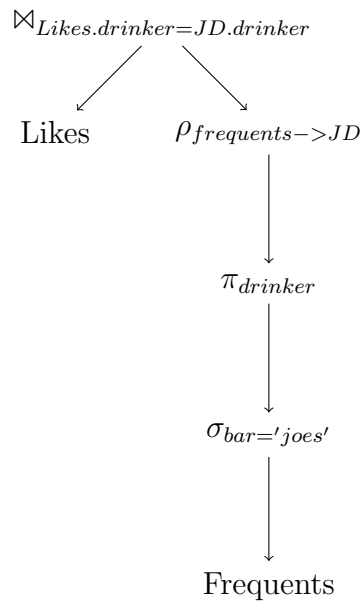
3.3 Subqueries

A parenthesized SELECT-FROM-WHERE statement

```

SELECT beer
FROM Likes, (SELECT drinker
              FROM Frequents                --create a table of drinkers who frequent joes
              WHERE bar = 'Joes') AS JD     --and call that table JD
WHERE Likes.Drinker = JD.drinker --now operate on that table

```



3.3.1 Query + Subquery Solution

Get the the bars where Miller is sold for the same price as bud at Joe's.

```

SELECT bar
FROM Sells
WHERE beer = 'Miller' AND
      price = (SELECT price --get the price from a subquery
              FROM Sells
              WHERE bar = 'Joes')

```

```
AND beer = 'Bud'  
LIMIT 1);      --Make sure you only get a single value
```

There is a runtime error if the subquery returns more than one value (hence the use of LIMIT).

3.3.2 IN

Get the name and manufacturer of each beer that Fred likes.

```
SELECT *  
FROM Beers  
WHERE name IN (SELECT beer  
FROM Likes  
WHERE drinker = 'Fred');
```

3.3.3 EXISTS

Correlated subquery: Walk one-by-one through tuples. At each tuple, perform the WHERE check. b1 refers to the tuple being checked.

```
SELECT name  
FROM Beers b1  
WHERE NOT EXISTS ( --only true if subquery returns empty set  
SELECT *      --get manufacturers who have only one beer  
FROM Beers  
WHERE manf = b1.manf AND  
name <> b1.name); --returns a tuple (manf, beer)
```

EXISTS looks at a table and finds out if the table returns any values. Either returns \emptyset or a table of results. **NOT EXISTS**(query) returns true only if the result of query is \emptyset . Internal query returns beers where the manufacturer is 'Bud' and the name is not 'Lite'.

3.3.4 ANY

$x = ANY(subquery)$ checks if x equals at least one tuple in the subquery.

3.3.5 ALL

$x <> ALL(subquery)$ is true iff x is not equal to any tuple in the subquery.

3.3.6 DISTINCT

Remove duplicates

SELECT x,y

3.4 Union, Intersect, Difference

Set operators. Require that arguments have same schema.

4 Bag Semantics

Just add ALL to any of the set operators e.g. UNION ALL.

δ remove duplicate tuples

τ sort tuples

γ grouping and aggregation

Sorting

$R1 := \tau_L(R2)$

L is a list of attributes to list by, listed in order of priority.

4.1 Grouping

$R1 := \gamma_L(R2)$

A	B	C
1	2	3
4	5	6
1	2	5

$\gamma_{A,B,AVG(C) \rightarrow X}(R) = ??$


First, group R by A and B.

A	B	C
1	2	3
1	2	5
4	5	6

Then, average C within groups

A	B	C
1	2	4
4	5	6

4.2 Outerjoin



4.2.1 JOIN Modifiers

R OUTER JOIN S is the core of an outerjoin. Can be modified by

- NATURAL before JOIN

4.3 Aggregations

Not operators of relational algebra, apply to an entire column.

Builtins: SUM, AVG, COUNT, MIN, MAX

Can define your own in postgres.

NOTE: NULL's are ignored in Aggregations.

4.3.1 Return a table with the single column AVG(Price).

```
SELECT AVG(price)
FROM Sells
WHERE beer = 'Bud';
```

4.3.2 Return number of distinct prices for bud beers

```
SELECT COUNT(DISTINCT price)
FROM Sells
WHERE beer = 'Bud';
```

4.4 Grouping

Follow SELECT-FROM-WHERE by GROUP-BY and a number of attributes

4.4.1 Get the Average price for each beer

```
SELECT Beer, AVG(price)
FROM Sells
GROUP BY beer;
```

Anything in GROUP BY must also be in SELECT.

In relational algebra: $\gamma_{beer, AVG(price)}$

Result:

beer	AVG(price)
bud	2.33
coors	4.00
...	...

4.4.2 Average Age by Gender

```
SELECT gender, AVG(age)
FROM People
GROUP BY age;
```

Result:

gender	AVG(age)
M	24
F	23

4.4.3 Illegal Example

```
SELECT bar, MIN(price)
FROM Sells
GROUP BY beer = 'Bud';
```

4.5 Having

Get only certain rows based on a condition.

4.5.1 Get average age only for genders where it is greater than 24

```
SELECT gender, AVG(age)
FROM People
GROUP BY gender
HAVING AVG(age) > 23
```

Result:

gender	AVG(age)
M	24